

Fake News Detection

A production-oriented **Fake News Detection** system for automatic classification of news articles based on textual content.

The project implements a complete, reproducible **machine learning pipeline** for NLP tasks, including data preprocessing, feature extraction, model training, evaluation, and result visualization. It is designed with clean architecture, detailed logging, and scalability in mind.



Features

- End-to-end ML pipeline: data loading → training → evaluation
 - Text vectorization using **TF-IDF**
 - Linear text classification with **SGDClassifier (Passive-Aggressive mode)**
 - Robust evaluation with accuracy and classification report
 - Visual analytics:
 - Confusion Matrix
 - Class distribution plots
 - Automatic saving of figures to `results/figures/`
 - Modular, production-ready project structure
 - Extensive logging for monitoring and debugging
-

Project Structure

```
fake_news_detection/
├── data/                      # Input datasets
│   └── fake_news.csv
├── results/                   # Generated visualizations
│   └── figures/
├── src/
│   ├── preprocessing.py        # Data loading & cleaning
│   ├── feature_extraction.py  # Train/test split & TF-IDF
│   ├── model.py                # Model training
│   ├── evaluation.py          # Metrics & visualizations
│   └── __init__.py
├── main.py                    # Application entry point
├── config.py                  # Project configuration
├── requirements.txt
├── pyproject.toml
└── README.md
```

Tech Stack

- Python 3.10+
 - scikit-learn
 - pandas, numpy
 - scipy
 - matplotlib, seaborn
-

Getting Started

Install dependencies

```
pip install -r requirements.txt
```

or with Poetry:

```
poetry install
```

Run the pipeline

```
from main import run_pipeline  
  
run_pipeline("data/fake_news.csv")
```

The pipeline performs:

1. Dataset loading and validation
 2. Stratified train/test split
 3. TF-IDF feature extraction
 4. Model training
 5. Model evaluation
 6. Visualization and artifact saving
-



ML Pipeline Overview

```
def run_pipeline(data_path: str):  
    df = create_dataframe_from_dataset(data_path)  
  
    x_train, x_test, y_train, y_test = split_dataset(  
        df, target_column="label", test_size=0.2  
    )
```

```

vectorizer = create_tfidf_vectorizer(max_df=0.7)
tfidf_train, tfidf_test, _ = vectorize_data(
    x_train, x_test, vectorizer
)

model = train_sgd_classifier(tfidf_train, y_train)

y_pred = model.predict(tfidf_test)
accuracy = evaluate_model(model, tfidf_test, y_test)

plot_confusion_matrix(y_test, y_pred)
plot_class_distribution(y_train)

return model, vectorizer

```

Model Performance

Typical results on the provided dataset:

- **Accuracy:** ~94%
- Balanced precision and recall across classes
- Stable convergence with early stopping

Generated artifacts:

- `results/figures/confusion_matrix_fake_news.png`
- `results/figures/class_distribution_train.png`

Model Choice

The project uses **SGDClassifier** configured to emulate the **Passive-Aggressive (PA-I)** algorithm:

```

SGDClassifier(
    loss="hinge",
    learning_rate="pa1",
    penalty=None,
    early_stopping=True
)

```

This approach provides:

- High efficiency on large, sparse text data
- Fast convergence
- Strong performance for binary text classification
- Compatibility with modern versions of scikit-learn

Roadmap & Extensions

Potential extensions for production or research use:

- Hyperparameter optimization (GridSearchCV)
 - n-gram feature enrichment
 - Full sklearn Pipeline integration
 - Model persistence (`joblib`)
 - REST API (FastAPI)
 - Real-time inference service
 - Monitoring and drift detection
-

License

MIT License